

Hadoop Put Command:

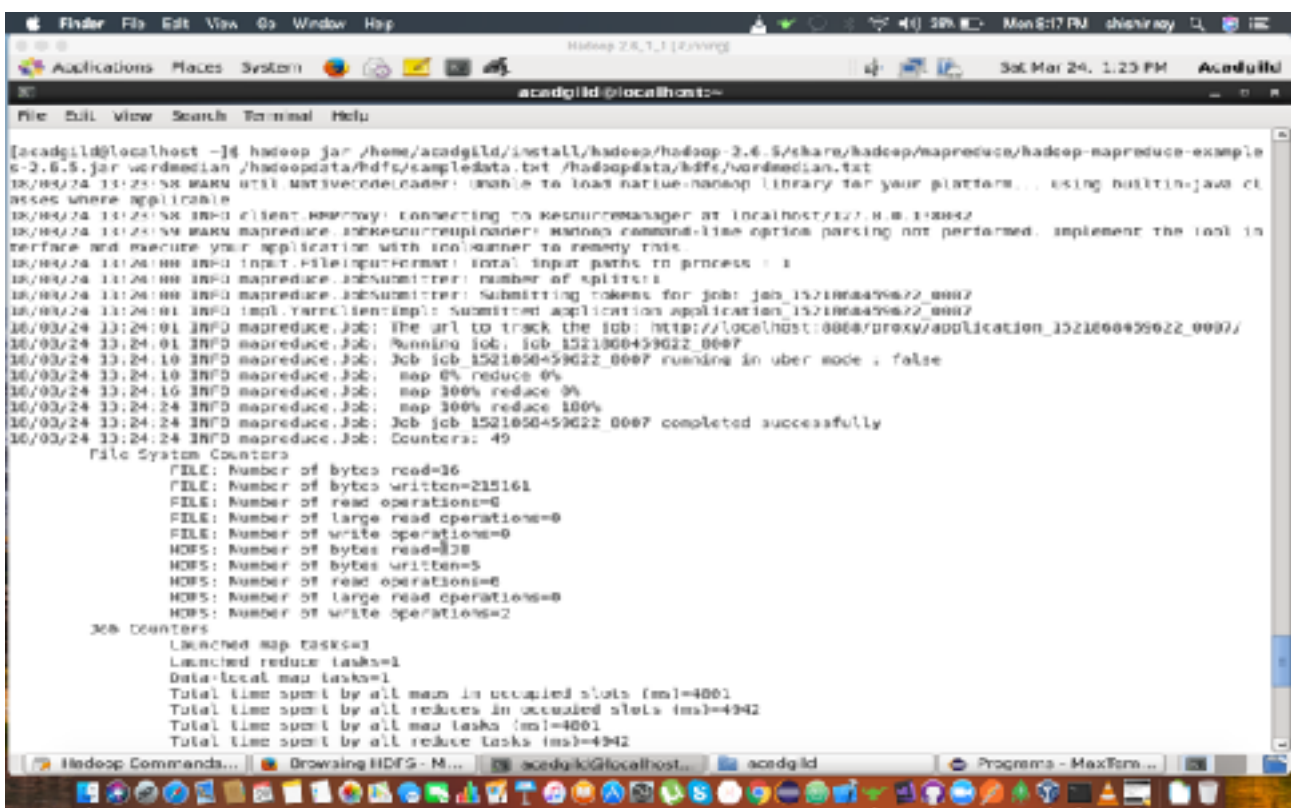
It is used to copy the file from local file to hadoop file system. So i have copied the sample data.txt for Execute WordMedian ,WordMean ,WordStandardDeviation programs using hadoop-mapreduce-examples-2.6.5.jar

For Example:

```
hadoop fs -put /home/acadgild/sampleddata.txt /hadoopdata/hdfs/sampleddata.txt
```

Executing hadoop-mapreduce-examples-2.6.5.jar for WordMedian data using below commands

```
hadoop jar hadoop-mapreduce-examples-2.6.5.jar /hadoopdata/hdfs/sampleddata.txt /hadoopdata/hdfs/wordmedian.txt
```



```
[acadgild@localhost ~]$ hadoop jar /home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.6.5.jar wordmedian /hadoopdata/hdfs/sampleddata.txt /hadoopdata/hdfs/wordmedian.txt
18/03/24 13:23:58 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/03/24 13:23:58 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8042
18/03/24 13:23:58 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the tool in
terface and execute your application with ToolRunner to remedy this.
18/03/24 13:24:00 INFO input.FileInputFormat: total input paths to process : 1
18/03/24 13:24:00 INFO mapreduce.JobSubmitter: number of splits:1
18/03/24 13:24:00 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1521866459622_0007
18/03/24 13:24:01 INFO impl.YarnClientImpl: Submitted application application_1521866459622_0007
18/03/24 13:24:01 INFO mapreduce.Job: The url to track the job: http://localhost:8080/proxy/application_1521866459622_0007/
18/03/24 13:24:10 INFO mapreduce.Job: Job job_1521866459622_0007 running in uber mode : false
18/03/24 13:24:10 INFO mapreduce.Job:  map 0% reduce 0%
18/03/24 13:24:16 INFO mapreduce.Job:  map 100% reduce 0%
18/03/24 13:24:24 INFO mapreduce.Job:  map 100% reduce 100%
18/03/24 13:24:24 INFO mapreduce.Job: Job job_1521866459622_0007 completed successfully
18/03/24 13:24:24 INFO mapreduce.Job: Counters: 49

File System Counters
  FILE: Number of bytes read=36
  FILE: Number of bytes written=215161
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=328
  HDFS: Number of bytes written=5
  HDFS: Number of read operations=6
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2

Job Counters
  Launched map tasks=1
  Launched reduce tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=4861
  Total time spent by all reduces in occupied slots (ms)=4942
  Total time spent by all map tasks (ms)=4861
  Total time spent by all reduce tasks (ms)=4942
```

```
Finder  File  Edit  View  Go  Window  Help
Hadoop 2.8.1.1 [Running]
Applications  Places  System
accdgild@localhost:~$
File  Edit  View  Search  Terminal  Help
Total score-milliseconds taken by all map tasks=1001
Total score-milliseconds taken by all reduce tasks=1042
Total megabyte-milliseconds taken by all map tasks=1007034
Total megabyte-milliseconds taken by all reduce tasks=1000000
Map-Reduce Framework
  map input records=1
  map output records=1
  map output bytes=1
  map output materialized bytes=1
  input split bytes=1
  combine input records=1
  combine output records=1
  Reduce input groups=1
  Reduce shuffle bytes=1
  Reduce input records=1
  Reduce output records=1
  Spilled Records=2
  Shuffled Maps=1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=125
  CPU time spent (ms)=1270
  Physical memory (bytes) snapshot=294506496
  Virtual memory (bytes) snapshot=4118245276
  Total committed heap usage (bytes)=170034480
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  ID_ERROR=0
  LENGTH=0
  MAP=0
  REDUCE=0
File Input Format Counters
  Bytes Read=21
File Output Format Counters
  Bytes Written=0
The session is: 20
You have new mail in /var/spool/mail/accdgild
accdgild@localhost ~$
```

Finder File Edit View Go Window Help
Hadoop 2.8.1.1 [Running]
Applications Places System
Browsing HDFS - Mozilla Firefox
Browsing HDFS All Applications
localhost:50070/explore.html#/hadoopdata/hdfs/wordman.txt Search
Hadoop Overview Datanodes Snapshot Startup Progress Utilities

Browse Directory

hadoopdata/hdfs/wordman.txt Go!

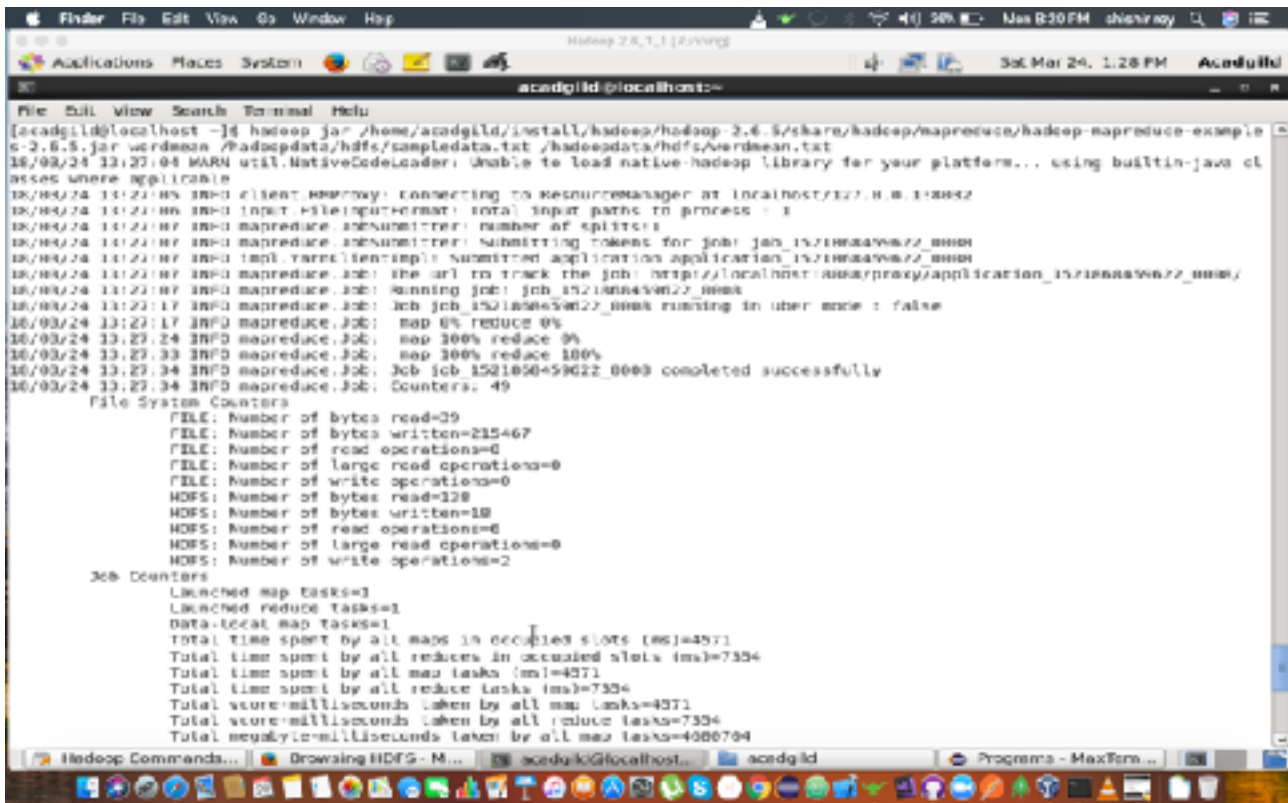
Permission	Owner	Group	Size	Replication	Block Size	Name
-rw-r--r--	accdgild	supergroup	0 B	1	128 MB	_SUCCESS
-rw-r--r--	accdgild	supergroup	10 B	1	128 MB	part-00000

Hadoop, 2016.

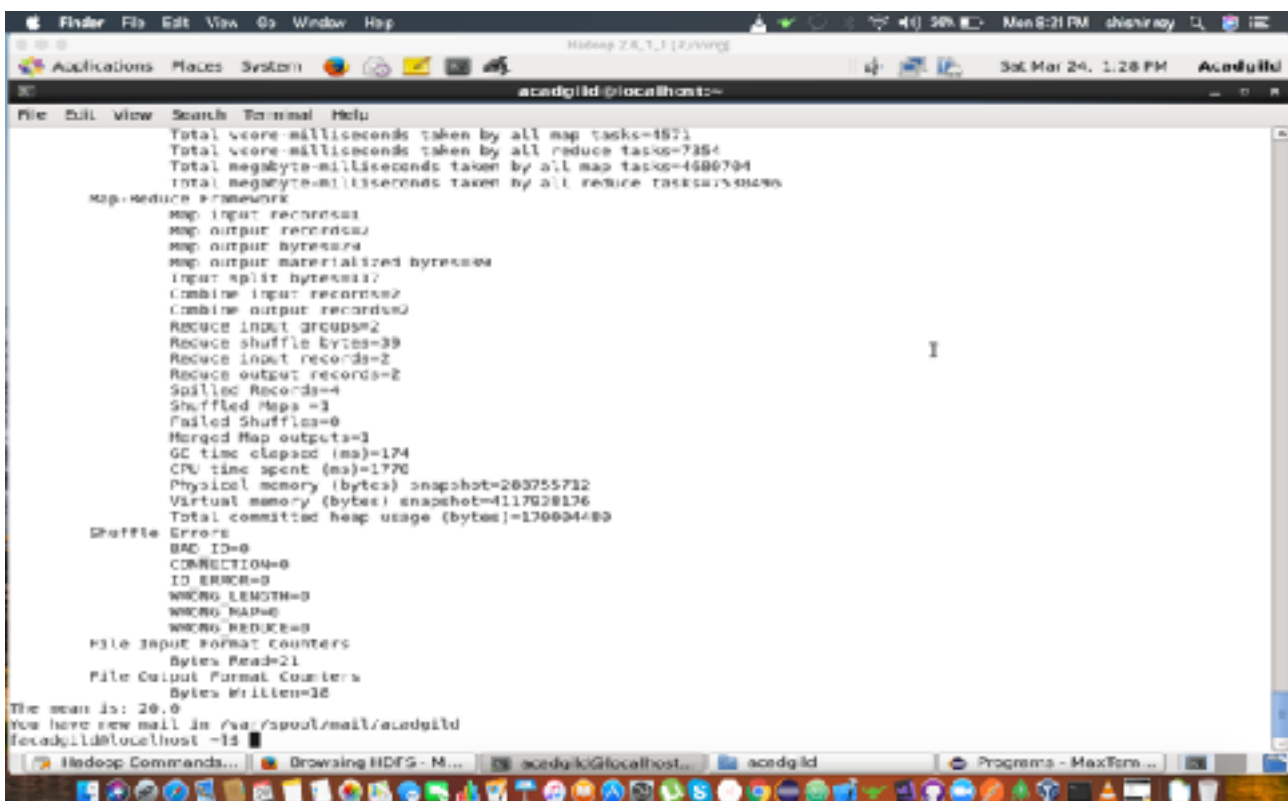
Hadoop Commands... Browsing HDFS - M... accdgild@localhost... accdgild Programs - MaxTerm...

Executing hadoop-mapreduce-examples-2.6.5.jar for WordMean data using below commands

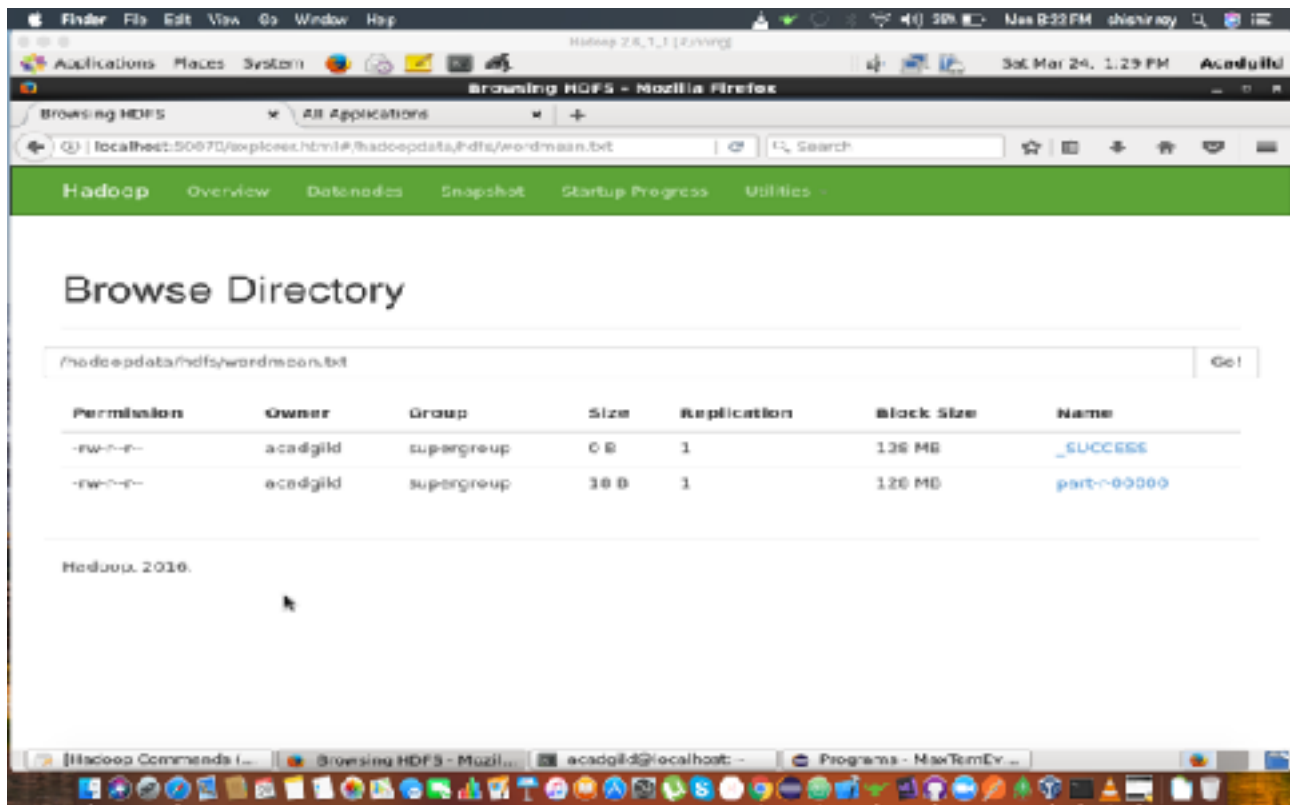
`hadoop jar hadoop-mapreduce-examples-2.6.5.jar /hadoopdata/hdfs/sampled.txt /hadoopdata/hdfs/wordmean.txt`



```
acsdgild@localhost:~$ hadoop jar /home/acsdgild/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.6.5.jar wordmean /hadoopdata/hdfs/sampled.txt /hadoopdata/hdfs/wordmean.txt
16/03/24 13:27:04 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
16/03/24 13:27:05 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
16/03/24 13:27:06 INFO input.FileInputFormat: total input paths to process : 1
16/03/24 13:27:07 INFO mapreduce.JobSubmitter: number of splits:1
16/03/24 13:27:07 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1521856459622_0000
16/03/24 13:27:07 INFO impl.YarnClientImpl: Submitted application application_1521856459622_0000
16/03/24 13:27:07 INFO mapreduce.Job: the url to track the job: http://localhost:8080/proxy/application_1521856459622_0000/
16/03/24 13:27:17 INFO mapreduce.Job: Job job_1521856459622_0000 running in uber mode : false
16/03/24 13:27:17 INFO mapreduce.Job:  map 0% reduce 0%
16/03/24 13:27:24 INFO mapreduce.Job:  map 100% reduce 0%
16/03/24 13:27:33 INFO mapreduce.Job:  map 100% reduce 100%
16/03/24 13:27:34 INFO mapreduce.Job: Job job_1521856459622_0000 completed successfully
16/03/24 13:27:34 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=39
    FILE: Number of bytes written=215467
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=128
    HDFS: Number of bytes written=18
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=4571
    Total time spent by all reducers in occupied slots (ms)=7354
    Total time spent by all map tasks (ms)=4571
    Total time spent by all reduce tasks (ms)=7354
    Total user-milliseconds taken by all map tasks=4571
    Total user-milliseconds taken by all reduce tasks=7354
    Total wall-clock-time taken by all map tasks=4580794
    Total wall-clock-time taken by all reduce tasks=7354
```

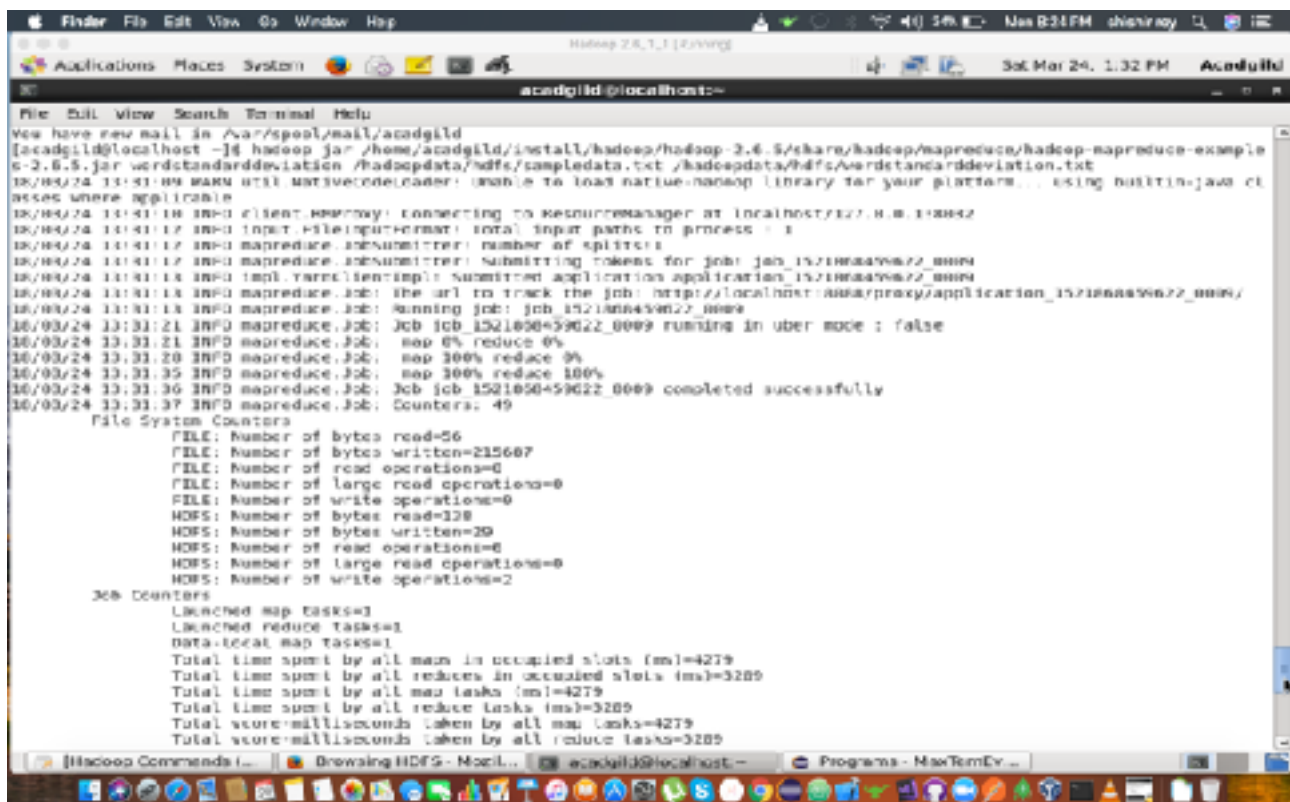


```
acsdgild@localhost:~$ hadoop jar /home/acsdgild/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.6.5.jar wordmean /hadoopdata/hdfs/sampled.txt /hadoopdata/hdfs/wordmean.txt
16/03/24 13:27:34 INFO mapreduce.Job: Job job_1521856459622_0000 completed successfully
16/03/24 13:27:34 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=39
    FILE: Number of bytes written=215467
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=128
    HDFS: Number of bytes written=18
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=4571
    Total time spent by all reducers in occupied slots (ms)=7354
    Total time spent by all map tasks (ms)=4571
    Total time spent by all reduce tasks (ms)=7354
    Total user-milliseconds taken by all map tasks=4571
    Total user-milliseconds taken by all reduce tasks=7354
    Total wall-clock-time taken by all map tasks=4580794
    Total wall-clock-time taken by all reduce tasks=7354
  Map-Reduce Framework
    Map input records=1
    Map output records=1
    Map output bytes=18
    Map output materialized bytes=18
    Input split bytes=128
    Combine input records=2
    Combine output records=2
    Reduce input groups=2
    Reduce shuffle bytes=39
    Reduce input records=2
    Reduce output records=2
    Spilled Records=4
    Shuffled Maps=1
    Failed Shuffles=0
    Merged Map outputs=1
    GC time elapsed (ms)=174
    CPU time spent (ms)=1770
    Physical memory (bytes) snapshot=263755732
    Virtual memory (bytes) snapshot=4117628136
    Total committed heap usage (bytes)=176934480
  Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    ID_ERROR=0
    LENGTH=0
    MAP=0
    REDUCE=0
  File Input Format Counters
    Bytes Read=21
  File Output Format Counters
    Bytes Written=18
The mean is: 20.0
You have new mail in /var/spool/mail/acsdgild
acsdgild@localhost:~$
```



Executing hadoop-mapreduce-examples-2.6.5.jar for WordStandardDeviation data using below commands

`hadoop jar hadoop-mapreduce-examples-2.6.5.jar /hadoopdata/hdfs/sampledata.txt /hadoopdata/hdfs/wordstandarddeviation.txt`



```
acsdgild@localhost:~$ cat /tmp/hadoop-ecsdgild/_logs/20160324-133213/hadoop-ecsdgild-20160324-133213-stdout.txt
Total score-milliseconds taken by all map tasks=4279
Total score-milliseconds taken by all reduce tasks=5286
Total megabyte-milliseconds taken by all map tasks=4381696
Total megabyte-milliseconds taken by all reduce tasks=5348816
Map-Reduce Framework
  Map input records=3
  Map output records=3
  Map output bytes=3
  Map output materialized bytes=3
  Input split bytes=3
  Combine input records=3
  Combine output records=3
  Reduce input groups=3
  Reduce shuffle bytes=59
  Reduce input records=3
  Reduce output records=3
  Spilled Records=6
  Shuffled Maps=3
  Failed Shuffles=0
  Merged Map outputs=3
  GC time elapsed (ms)=125
  CPU time spent (ms)=1296
  Physical memory (bytes) snapshot=265261728
  Virtual memory (bytes) snapshot=4118245326
  Total committed heap usage (bytes)=176934489
Shuffle
  Errors
    BAD_ID=0
    CONNECTION=0
    ID_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=21
File Output Format Counters
  Bytes Written=25
The standard deviation is: 0.3
You have new mail in /var/spool/mail/acsdgild
ecsdgild@localhost ~$
```

